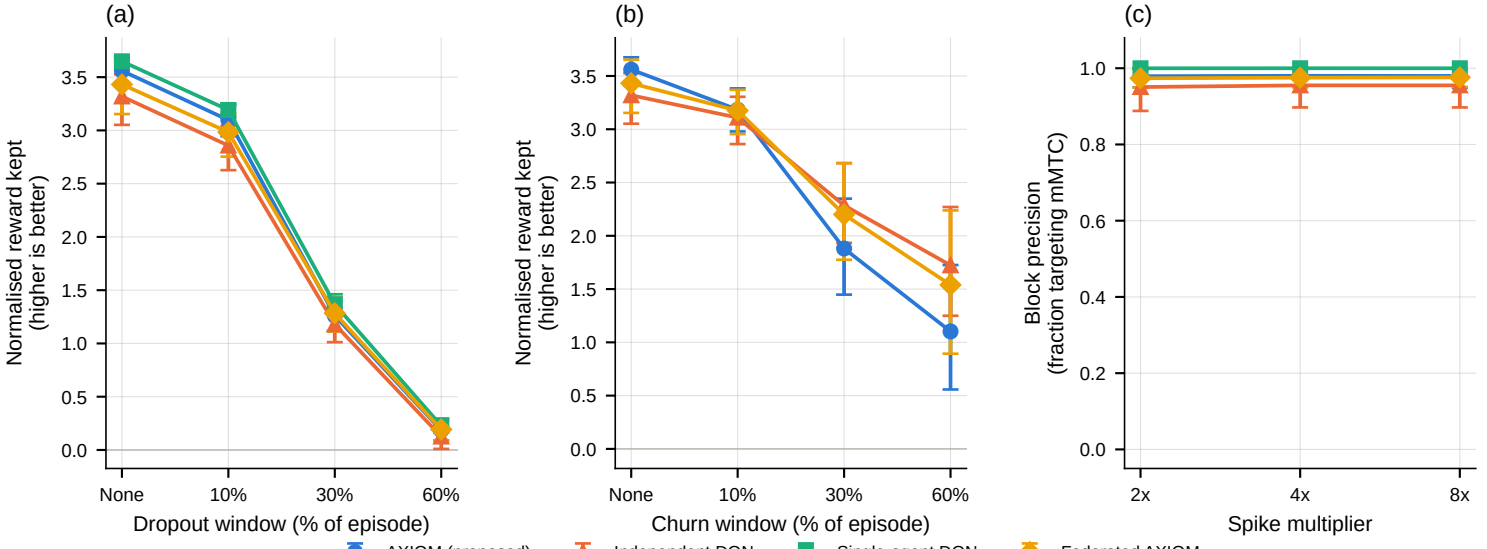
(a)
Normalised reward kept
(higher is better)
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
None
10%
30%
60%
Dropout window (% of episode)
(b)
Normalised reward kept
(higher is better)
3.5
3.0
2.5
2.0
1.5
1.0
0.5
0.0
None
10%
30%
60%
Churn window (% of episode)
(c)
Block precision
(fraction targeting mMTC)
1.0
0.8
0.6
0.4
0.2
0.0
2x
4x
8x
Spike multiplier
AXIOM (proposed)
Independent DQN
Single-agent DQN
Federated AXIOM