a

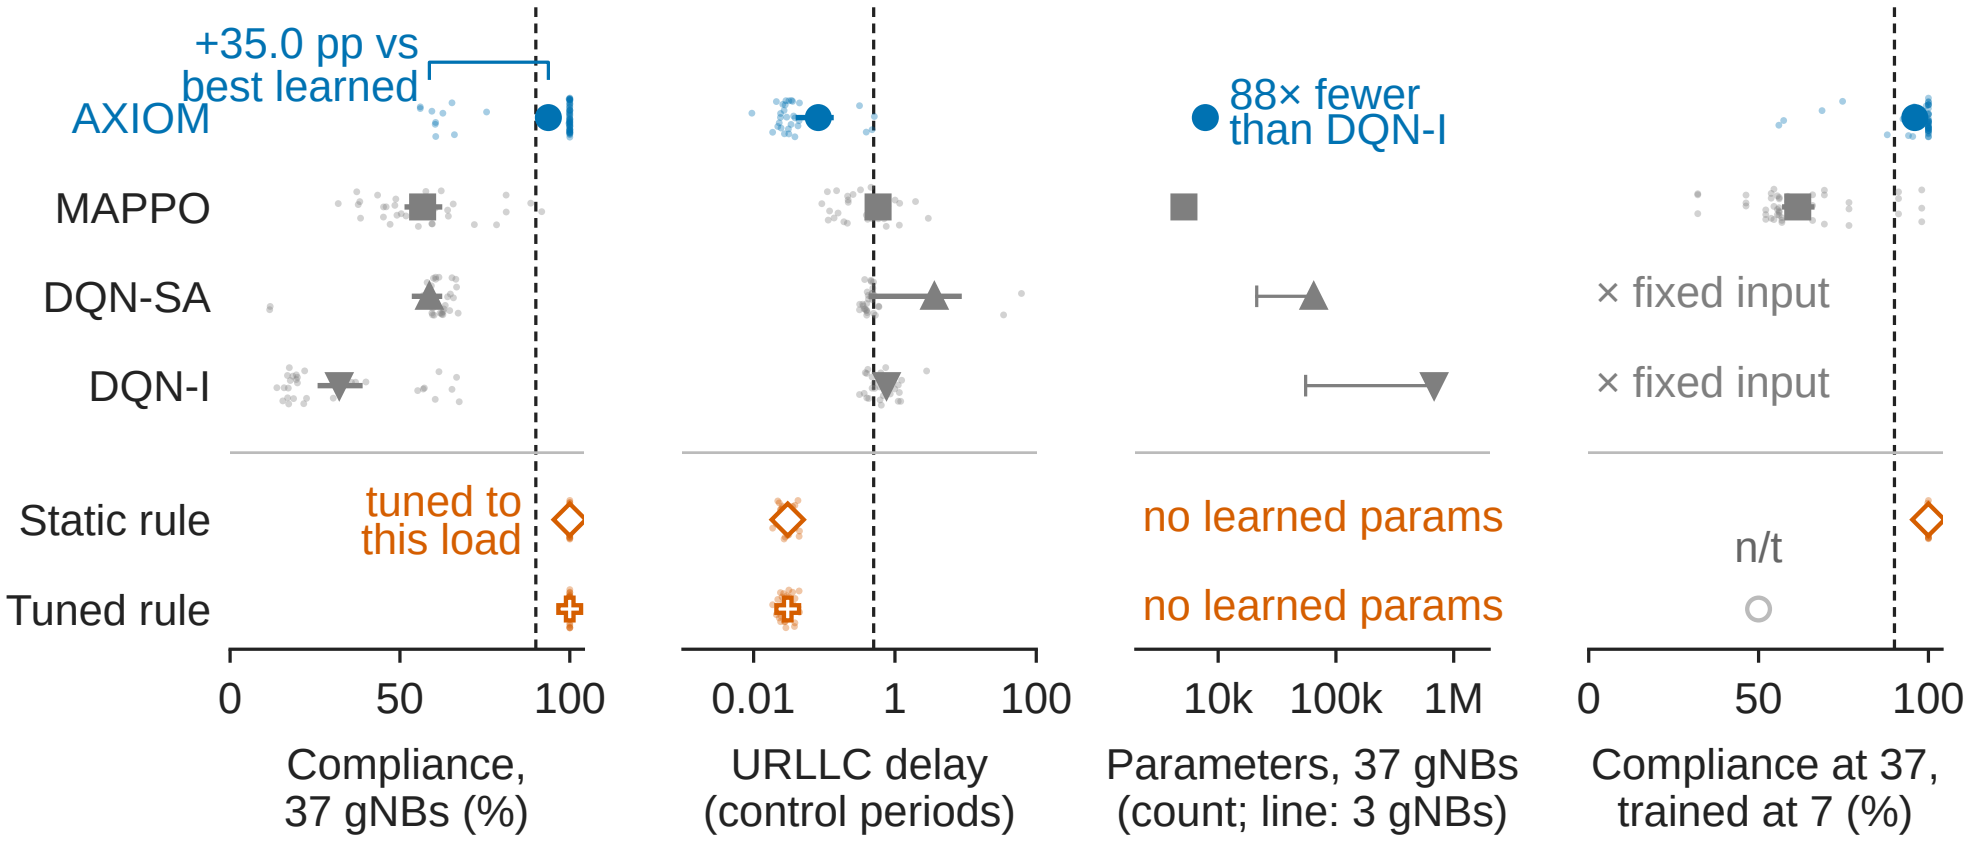

b

	AXIOM	MAPPO	DQN-SA	DQN-I	Static rule	Tuned rule
Transfers without retraining	✓ 95–96%	◐ 61.5%	× fixed input	× fixed input	✓ size-free	✓ 7→19, Env B
Beats tuned rule past its load	✓ +7.5/step	× 0 of 9 wins	◐ 4 of 9 wins	◐ 3 of 9 wins	✓ +9.4/step	reference
Best under steady load	× −3.8/step	× −25.0/step	× −9.4/step	× −12.9/step	✓ best, Env A	✓ best, A and B
Private training	✓ FedAvg	× pooled data	× pooled data	✓ local only	✓ no model	✓ no model