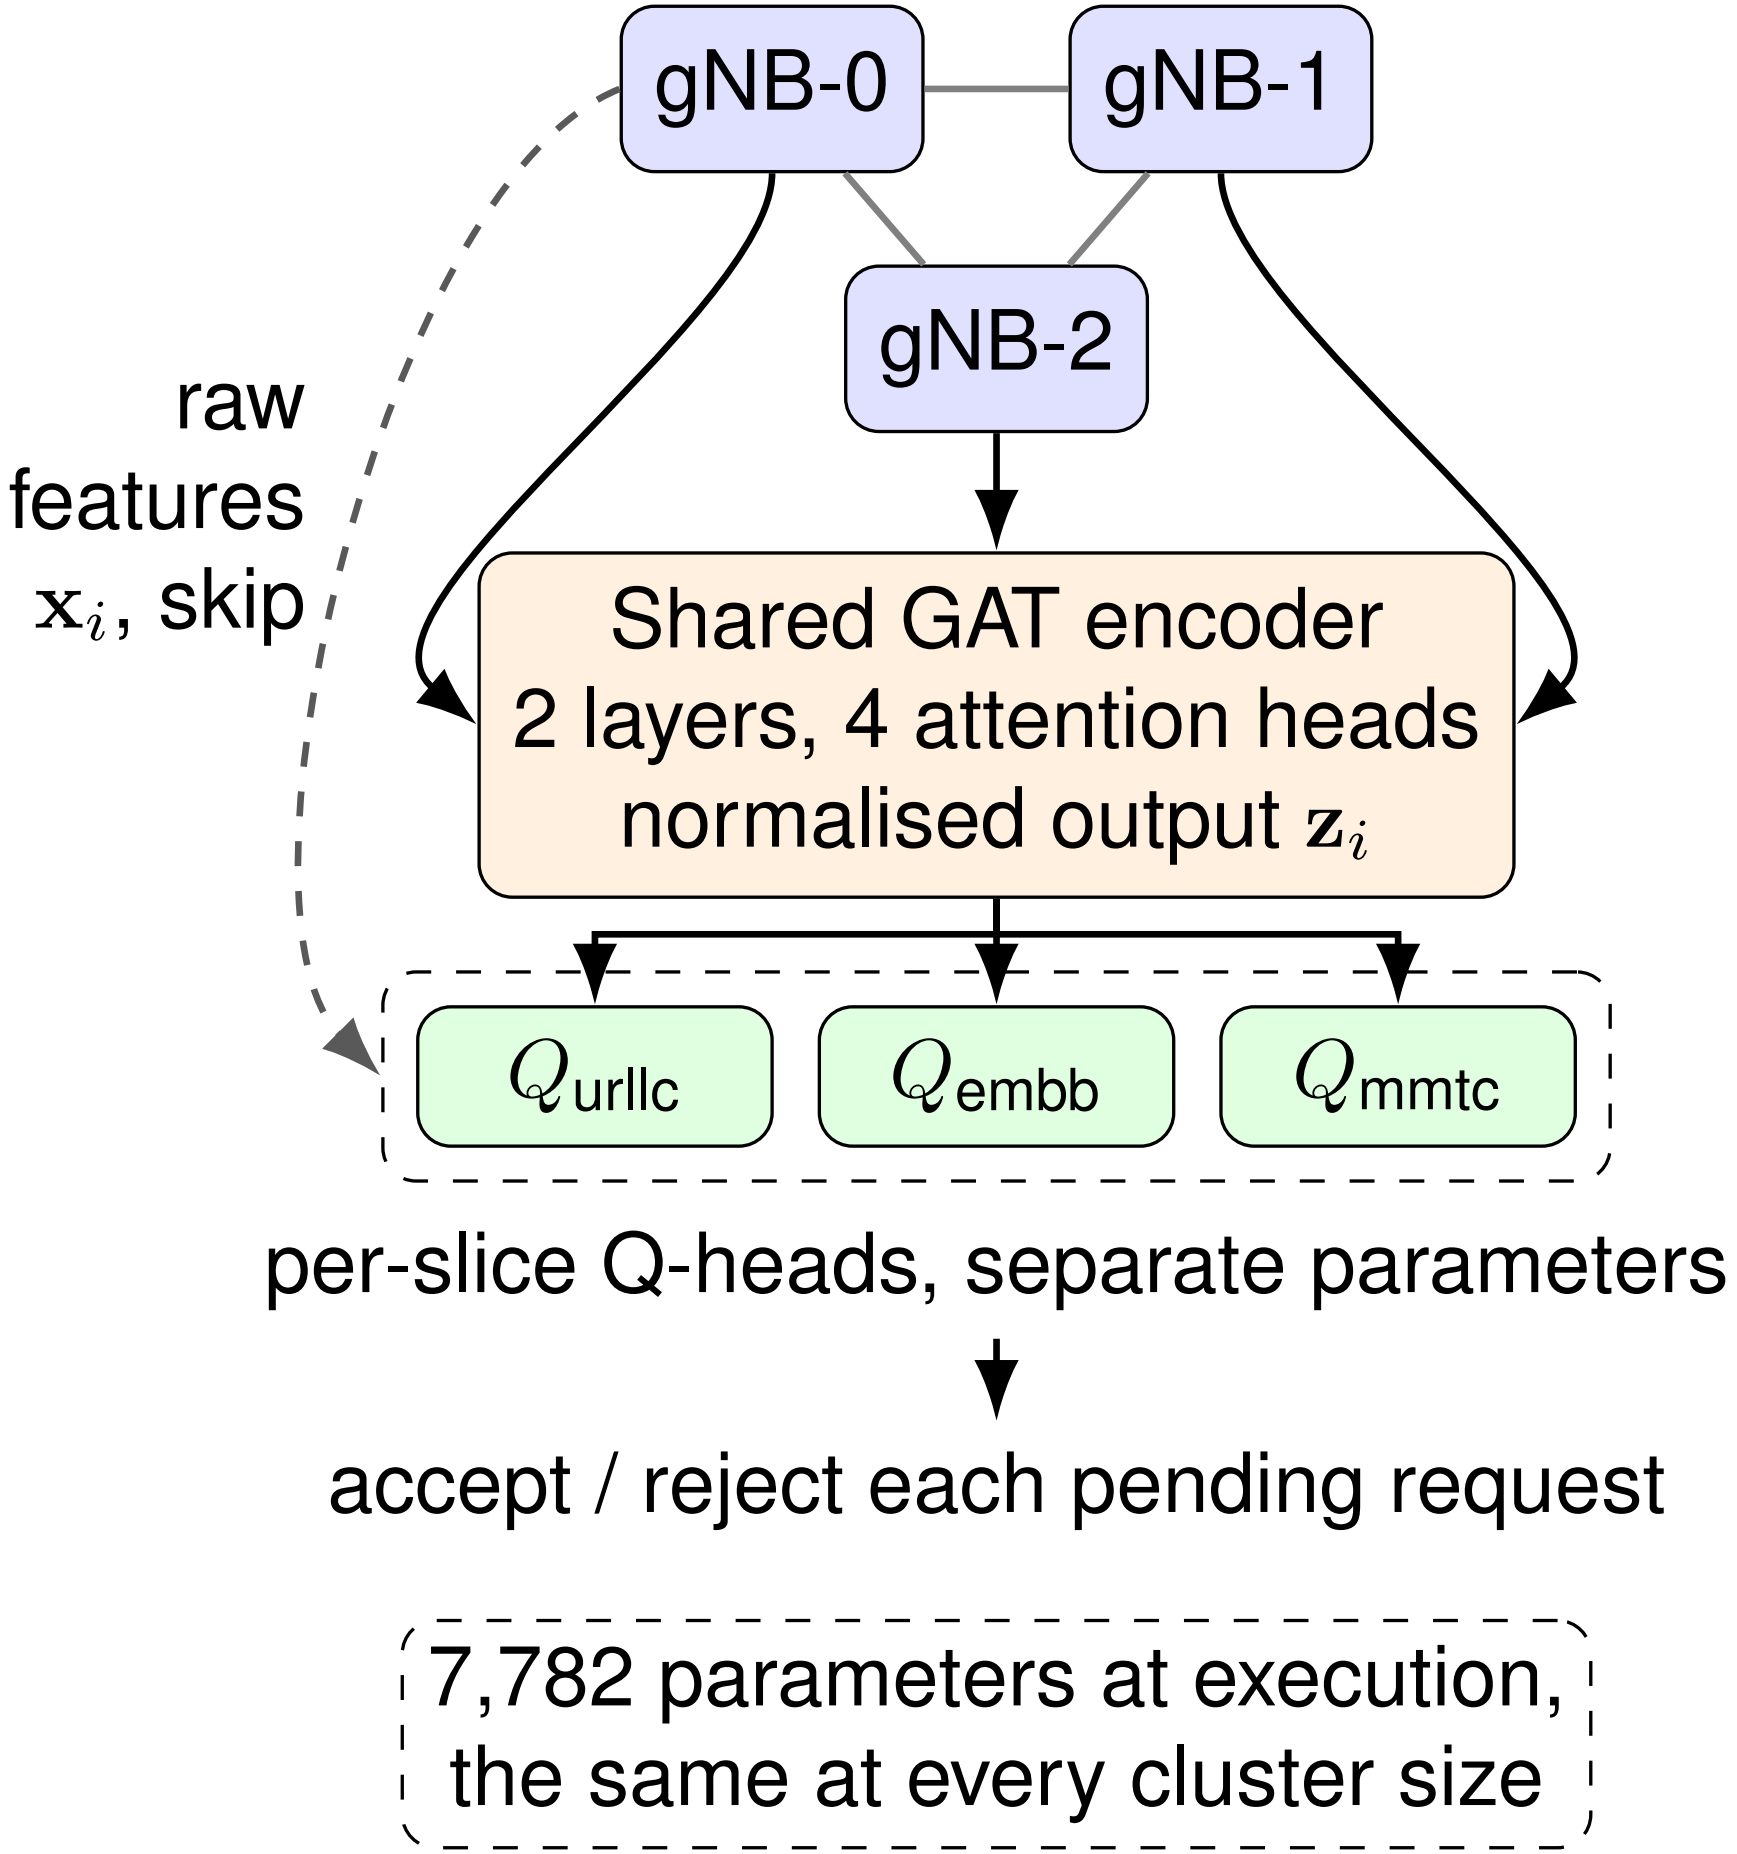
gNB-0
gNB-1
gNB-2
raw features $\mathbf{x}_i$, skip
Shared GAT encoder
2 layers, 4 attention heads
normalised output $\mathbf{z}_i$
Q_{urllc}
Q_{embb}
Q_{mmtc}
per-slice Q-heads, separate parameters
accept / reject each pending request
7,782 parameters at execution,
the same at every cluster size