

From Architecture to Evidence: Realising a Graph-Attention, Multi-Agent, Federated Slice Admission Controller Under Correctness-Aware Evaluation

Manoj Prasad Kunasegran¹, Wai Leong Pang^{1*}, Swee King Phang¹

^{1*}School of Engineering, Taylor’s University, Subang Jaya, Malaysia.

*Corresponding author(s). E-mail(s): TODO@taylors.edu.my;

Abstract

This paper delivers a working realisation of a graph-attention, multi-agent, federated slice admission controller, GAT-CTDE, and treats getting its evaluation right as a first-class contribution alongside the architecture itself. GAT-CTDE couples a shared Graph Attention encoder over the gNB cluster topology to decentralised per-slice Q-heads, trained centrally and executed independently per gNB, then extended to federated training with differential privacy and evaluated under mid-episode disruption. On the metric that reflects the actual training objective, it decisively beats an independent-learner ablation that isolates the same coordination contribution, though its edge over a single-agent baseline is small and not statistically decisive at this sample size, an honest limit we report rather than obscure. Scaled from 3 to 19 gNBs under a disclosed, deliberately harder-than-live stress regime, GAT-CTDE resists a training collapse that becomes total for the single-agent baseline, converging, across 64 independently-drawn seeds, on a residual 35.4% [25.5%, 45.8%] collapse rate, a pattern a properly powered resample shows is already present at 7 gNBs, not only at large cluster size. Federating the architecture costs nothing measurable in reward, and differential privacy degrades decision precision only past a real threshold rather than smoothly; under gNB dropout, demand spikes, and agent churn, its cost accelerates with severity like every arm tested. Two safeguards make these results trustworthy rather than merely favourable: a correctness-aware metric pair that caught a real training collapse a standard SLA-compliance metric scored no worse than correct behaviour, replicated across 44 of 49 independent checks, and a reproduction-versus-replication discipline that caught and retracted one architectural overclaim before publication. A live single-gNB deployment confirms the offline-trained decision logic executes correctly on real OAI hardware end to end, and surfaces a genuine cross-layer integrity finding: a scheduler-configuration bug had silently manufactured the appearance of contention in earlier live data,

and once fixed at its root, this rig cannot reach genuine per-slice scarcity at any load level yet tested, so the live result establishes deployability rather than admission-control behaviour under contention, which this paper tests instead in the disclosed offline stress regime.

Keywords: Network Slicing, Slice Admission Control, Multi-Agent Reinforcement Learning, Graph Attention Networks, Federated Learning, Open RAN

M. P. Kunasegran: orcid.org/0009-0005-3169-1873. W. L. Pang: orcid.org/0000-0001-8407-5648.
S. K. Phang: orcid.org/0000-0002-7877-8766.

1 Introduction

Network slicing lets a single physical radio access network serve heterogeneous service classes, namely Ultra-Reliable Low-Latency Communication (URLLC), enhanced Mobile Broadband (eMBB), and massive Machine-Type Communication (mMTC), under distinct per-slice Service Level Agreements (SLAs). Our earlier work [7, 8] established that a Deep Q-Network (DQN) admission controller can learn to enforce these SLAs under a single gNodeB (gNB), validated on a live OpenAirInterface (OAI) testbed. A companion submission [10] extends this single-gNB live validation to a QoE-aware reward variant, independently confirming the live E2 control loop’s feasibility (median round-trip latency 0.57 ms) for a policy under 20,000 parameters. A real multi-cell deployment does not admit requests at one gNB in isolation, since neighbouring cells share backhaul, spectrum planning, and a cluster-wide load-balance objective, making coordinated admission a genuinely multi-agent problem. This raises a direct question: does a shared, graph-structured representation of the cluster’s state help a set of per-gNB agents make better joint decisions than a single agent using the flattened joint state, or independent per-gNB agents with no shared representation at all?

This paper’s contribution is GAT-CTDE, a centralised-training, decentralised-execution multi-agent architecture built on a shared Graph Attention Network [19] encoder over the gNB topology feeding decentralised per-slice Q-heads, delivered together with the evaluation discipline needed to trust what it reports. Three gaps motivate this pairing, each named directly against our own systematic review of 67 core studies (Section 2): no reviewed study checks whether its own outcome metric can score a degenerate policy no worse than a correct one; none replicates a headline architectural finding on an independent seed sample before reporting it; and none tests whether a live testbed can actually reach the contention regime its offline claims are evaluated under. We close all three: a per-component realisation of the architecture (GAT encoder, MARL controller, FL coordinator, DP layer), each with its own measured, isolated contribution rather than one end-to-end number; a correctness-aware metric pair and reproduction-versus-replication discipline applied throughout rather than asserted once; and a live single-gNB deployment that tests, rather than assumes, whether this rig’s own achievable load ever produces genuine per-slice scarcity.

Extended to federated training with differential privacy and evaluated under mid-episode gNB dropout, demand spikes, and agent churn, we scale the architecture to cluster sizes $N \in \{7, 19\}$ under ring and hex-grid adjacency. It resists a training collapse that becomes total for a flattened-state single-agent baseline, converging on a residual 35.4% collapse rate across 64 independently-drawn seeds rather than either of two disagreeing small-sample estimates, a pattern already present at $N=7$ rather than only emerging at large cluster size. On the metric that reflects the actual training objective, the architecture decisively beats an independent-learner ablation that isolates the same coordination contribution, though its edge over the single-agent baseline specifically is small and not statistically decisive at this sample size. This is an honest limit we report rather than smooth over to fit a cleaner architecture story. Federating the architecture costs nothing measurable in reward, and differential privacy degrades decision precision only past a real threshold rather than a smooth trade-off; under disruption, its cost accelerates with severity like every arm tested, with a coordination-specific penalty at only the mildest severity tested.

Reporting these results honestly required two methodological safeguards we also contribute, since neither is standard practice in this literature. First, the standard episode-level SLA-compliance metric can score a degenerate, always-accept policy no worse than genuinely differentiated decision-making, because blocking a low-priority slice under congestion is the reward-optimal action, yet it costs compliance without ever recovering it. We introduce a correctness-aware metric pair (mean reward per step and block precision, both computed from data every run already logs) that exposes this gap, and use it to diagnose and fix a real training collapse our first encoder suffered on all 30 campaign seeds. Second, we establish a reproduction-versus-replication discipline. Every headline result is checked against an independent seed sample rather than only reproduced on the same seeds, a distinction that caught and retracted a real architectural overclaim under mid-episode disruption before it could be reported as a finding. We finally deploy the architecture’s single-gNB baseline on a real OAI testbed and treat what that deployment can and cannot establish as itself a finding: the offline-trained decision logic executes correctly on real hardware, and a cross-layer scheduler misconfiguration that had earlier manufactured the appearance of live contention is fixed at its root, after which this rig cannot exercise genuine per-slice scarcity at any load level it has been tested at, confining the disruption and scaling claims above to the disclosed offline stress regime where they are made.

2 Related Work

Our own systematic review [9] screened 6,286 records down to 67 core 2024–2026 studies on machine-learning-driven, user-centric distributed network slicing, organised around three enabling pillars, namely topology-aware graph representation, distributed multi-agent control, and privacy-preserving federation, and found individual pillars well developed but rarely combined under realistic constraints, and never with a validated live deployment. None of the works we reviewed, there or here, checks whether its own evaluation metric can mask a degenerate policy, or replicates a headline finding on an independent seed sample. This is the gap this paper fills.

The canonical result behind our own encoder choice is Shao *et al.* [16], who show a graph attention encoder strengthens inter-agent cooperation over plain DQN/A2C baselines in dense-cellular slicing. Closest to our admission-control setting, Ahmadi *et al.* [2] generalise coordinated multi-agent slicing across unseen topologies via graph attention, a topology-generalisation question distinct from, but complementary to, the evaluation-integrity question we ask on a fixed topology. We test the complementary direction directly in Section 7, asking whether correctness-aware metrics remain reliable as cluster size and adjacency sparsity change, rather than whether a policy generalises to an unseen topology. Sulaiman *et al.* [17] is the closest prior art to our own problem, multi-agent DRL jointly solving slicing and admission control across the substrate network; neither that work nor its topology-generalisation follow-up reports whether its own reward-optimal policy could be mistaken for a degenerate one, or checks a headline result against an independent seed sample. Tashman and Cherkaoui [18] and Fatehi *et al.* [5] both study robustness to a disruptive condition, as our disruption campaign (Section 6.3) does; neither checks whether the metric reporting that robustness is itself trustworthy, or replicates a robustness finding on an independent sample. This is exactly the gap our own disruption campaign’s independent-seed replication closes, and not hypothetically, since it is what overturned our own initial churn-immunity reading.

On the federated side, Zhang *et al.* [21] coordinate independent O-RAN xApps via federated DRL, the closest precedent to our own periodic- synchronisation motivation, but without a differential-privacy step or a correctness-aware metric to check whether federating changed decision quality. Yasin *et al.* [20] combine federated edge learning with differential privacy to secure the O-RAN RIC (the closest FL+DP precedent) but target a different threat model (securing the controller) rather than characterising a control policy’s own privacy-utility trade-off, which our σ -sweep (Section 6.2) reports. A small further set of works begins to combine two of these three pillars [3, 6, 14, 15], but, like every work surveyed here, evaluates exclusively on its own reported outcome metric, computed once, on one seed or run. This is the gap this paper fills, not a new point in the GAT/MARL/FL design space but a demonstration (on GAT-CTDE, extended to federation, privacy, and disruption) that checking a metric’s own trustworthiness and replicating a finding on an independent sample are both necessary, and that skipping either changes what a reader is entitled to believe the resulting numbers mean.

3 System Model and Problem Formulation

We consider a cluster of $N=3$ gNBs, each serving the same three slices (URLLC, eMBB, mMTC) under a shared per-gNB Physical Resource Block (PRB) budget B . An admission decision is realised as a PRB-ceiling adjustment (`slicing_control_m` min/max ratio), not a literal per-flow accept/reject, since the real OAI xApp control API only exposes ratio ceilings. No real multi-gNB physical topology is available to us, so we take the least assumption-laden choice and treat the cluster as fully connected, noting this explicitly as a modelling choice rather than a measured fact. At each control step, every gNB has zero or more pending admission requests, each tagged with a

slice identity, and an agent decides accept/reject for each independently, conditioned on the current cluster state; the discrete action space matches our single-gNB work exactly (`action_dim=2`), so any coordination benefit found here is attributable to the architecture, not a richer action space.

The reward combines a per-slice SLA-compliance term with a cluster-wide load-balance penalty, weighted by a per-slice `priority_weight` (URLLC highest, mMTC lowest) and a congestion coefficient tuned, in prior work, so that the reward-optimal policy under congestion is *differentiated shedding*: continue accepting URLLC, mostly accept eMBB, and reject mMTC. This is confirmed directly by Q-value inspection on a working baseline policy (Section 5), not assumed. Every result in this paper except Section 8’s bounded live anchor is obtained in an offline, closed-loop simulation environment deliberately oversubscribed to roughly 110% of one gNB’s PRB budget, reusing our existing multi-gNB configuration unmodified. A separate investigation found this offline environment’s demand process does not reproduce the live testbed’s SLA-margin distribution closely enough to support a live-rank prediction claim, even after recalibrating the demand model’s temporal structure. A live re-measurement closes this further: at every native load level this rig has been tested at, one and two UEs per slice alike, all three slices serve essentially 100% of offered demand with zero loss (Section 8), meaning this rig has never been shown to reach a load at which any slice’s real admission decision carries a genuine capacity consequence. We therefore use this offline environment explicitly as a deliberately harder, disclosed *stress environment* for a contention regime this live rig cannot currently produce, not as a miscalibrated approximation of a reachable live condition. Multi-gNB and disruption claims in this paper remain offline-only throughout; this rig cannot yet sustain a live multi-gNB test to check them against.

4 GAT-CTDE Method

4.1 Shared Encoder, Decentralised Q-Heads

Fig. 1(a) shows the architecture. At every control step, each gNB’s own per-slice [PRB-used-ratio, congestion-level, queue-length-norm] features form that node’s row in a joint node-feature matrix. A two-layer, four-head GAT [19], implemented directly (the small graph sizes here make a dense, masked-attention implementation as fast as a sparse one), produces one embedding per gNB. Centralisation comes from this encoder, since its parameters are updated by every agent’s loss jointly, so it sees the full cluster’s state during training. Decentralised execution comes from what happens next. At action-selection time, each agent consumes only its own node’s embedding row to decide its own pending requests.

An early version fed each embedding, concatenated with a one-hot slice context, into a single shared Q-head MLP. Under this design, mMTC’s true reward-optimal margin under congestion is small by the reward’s own calibration, while URLLC’s much larger priority weight produces proportionally larger gradients; because all three slices update the *same* shared parameters, URLLC’s abundant, large-magnitude gradients overwrite mMTC’s much smaller true signal before it registers. This is a standard shared-parameter/class-imbalance failure mode. We replace the single shared head




Fig. 1 GAT-CTDE architecture. (a) A shared Graph Attention encoder over the fully-connected 3-gNB topology produces one embedding per node; per-slice Q-heads, each with its own parameters, consume that node’s own embedding row at action-selection time (decentralised execution). (b) The federated training variant: each gNB is a client with a private copy of the same network, synchronised via periodic FedAvg with an optional per-client DP-SGD clip-and-noise step.

with *one small MLP per slice*, selected by the context one-hot where the head is applied (Fig. 1(a)). Each slice’s decision boundary now has its own parameters, so one slice’s gradients can no longer directly overwrite another’s.

Algorithm 1 states the training loop, in which the encoder is updated from every request in a sampled minibatch regardless of gNB or slice (centralised training), while action selection at rollout time consults only the acting agent’s own embedding row (decentralised execution). Every pending request in a step is its own independent temporal-difference sample, never summed or averaged across a step before the loss is computed, matching our existing single-agent DQN’s own granularity.

4.2 Federated Training with Differential Privacy

Fig. 1(b) shows the federated variant, in which each gNB becomes an FL client with a private copy of the same network, training only on its own requests. Transitions are partitioned by owning agent, so no client sees another client’s raw admission data, while its GAT encoder still consumes the full joint node-feature snapshot every arm observes. Every R local steps, the server averages all clients’ online weights (FedAvg [13]) and broadcasts the result back. We also implement FedProx [12] (a proximal term pulling local weights toward the last broadcast global model), motivated by an implicit, seeded per-gNB load multiplier this project’s own cluster-scaling investigation (Section 7) later revealed as genuine, uncontrolled client heterogeneity. Swept at $\mu \in \{0.01, 0.1, 1.0\}$ under both a homogeneous-load control and this heterogeneity, FedProx earns no measurable dividend at any tested strength. This traces to a concrete mechanism, since real per-round client drift within `local_steps_per_round` averages only 0.03% of the TD loss, too small for the proximal term to matter. A

Algorithm 1 GAT-CTDE: centralised training, decentralised execution

Require: shared GAT encoder f_θ , per-slice Q-heads $\{Q_{\phi_k}\}_{k=1}^K$, target networks $f_{\theta^-}, \{Q_{\phi_k^-}\}$, replay buffer $\mathcal{D}$, fully-connected adjacency A , discount γ , target-sync period C episodes

- 1: **for** episode = 1, 2, ... **do**
- 2: reset environment; observe joint node features $X_1 \in \mathbb{R}^{N \times d}$
- 3: **for** step $t = 1, \dots, T$ **do**
- 4: $E_t \leftarrow f_\theta(X_t, A)$ {one embedding row per gNB; centralised: sees every node}
- 5: **for** each pending request r at gNB i , slice k **do**
- 6: $a_r \leftarrow \arg \max_a Q_{\phi_k}(E_t[i])[a]$ w.p. $1-\epsilon$, else uniform random {decentralised: own row $E_t[i]$ only}
- 7: **end for**
- 8: $X_{t+1}, \rho_t, \text{done} \leftarrow \text{ENV.STEP}(\{a_r\})$
- 9: store $(X_t, X_{t+1}, \{(i, k, a_r)\}_r, \rho_t, \text{done})$ in $\mathcal{D}$
- 10: **if** $|\mathcal{D}| \geq \text{warmup size}$ **then**
- 11: sample minibatch $B \subset \mathcal{D}$
- 12: $E_B \leftarrow f_\theta(X_B, A)$; $E'_B \leftarrow f_{\theta^-}(X'_B, A)$ {target net, no grad}
- 13: **for** each request $(i, k, a) \in B$ **do**
- 14: $q \leftarrow Q_{\phi_k}(E_B[i])[a]$
- 15: $y \leftarrow \rho + \gamma \cdot (1 - \text{done}) \cdot \max_{a'} Q_{\phi_k^-}(E'_B[i])[a']$
- 16: **end for**
- 17: $\mathcal{L} \leftarrow \text{mean Huber loss over every request's } (q, y) \text{ pair}$ {per-request, never summed per step}
- 18: update $\theta, \{\phi_k\}$ by $\nabla \mathcal{L}$; clip grad-norm to 1.0
- 19: **end if**
- 20: **end for**
- 21: **if** episode mod $C = 0$ **then**
- 22: $\theta^- \leftarrow \theta$; $\phi_k^- \leftarrow \phi_k \ \forall k$
- 23: **end if**
- 24: **end for**

follow-up at $10\times$ the round length lets real drift grow nearly 50-fold (to a mean 1.44%) yet still produces eval-time behaviour bit-identical to FedAvg, ruling out round length as the limiting factor rather than leaving it untested.

Differential privacy is applied as a standard DP-SGD [1] step, in which every client's gradient is clipped to a fixed norm and, when enabled, perturbed with Gaussian noise scaled by a noise multiplier σ , before that client's own optimiser step. We report σ as the primary privacy parameter, alongside a formal (ϵ, δ) bound computed via zero-concentrated differential privacy (zCDP) [4] composition of the Gaussian mechanism ($\rho = k/(2\sigma^2)$ over k per-client releases, $\epsilon(\delta) = \rho + 2\sqrt{\rho \ln(1/\delta)}$). This is a conservative, non-subsampled bound, reported alongside σ rather than in place of it. Algorithm 2 states one federated round precisely; clipping is applied identically whether or not noise is added, so a $\sigma=0$ run and a $\sigma>0$ run differ only in whether noise is injected. This is the isolation the privacy-cost comparison (Section 6.2) depends on.

Algorithm 2 Federated GAT-CTDE round with DP-SGD

Require: N clients, each with local network $(\theta_i, \{\phi_{i,k}\})$; global snapshot $\bar{\theta}$; local steps per round R ; clip norm C_{dp} ; noise multiplier σ ; FedProx weight μ ($= 0$ unless noted)

```
1: for round = 1, 2, ... do
2:   for local step = 1, ...,  $R$ , each client  $i$  independently do
3:     sample minibatch  $B_i$  of  $i$ 's own requests from  $i$ 's local buffer
4:      $E \leftarrow f_{\theta_i}(X_{B_i}, A)$  {full joint state; client's OWN local weights only}
5:      $\mathcal{L}_i \leftarrow$  per-request Huber TD loss (Alg. 1, lines 11–13) +  $\frac{\mu}{2} \|\theta_i - \bar{\theta}\|^2$  {FedProx term}
6:      $g_i \leftarrow \nabla_{\theta_i, \phi_i} \mathcal{L}_i$ 
7:     if  $\sigma > 0$  then
8:        $g_i \leftarrow \text{clip}(g_i, C_{\text{dp}})$ ;  $g_i \leftarrow g_i + \mathcal{N}(0, (\sigma C_{\text{dp}})^2 I)$  {DP-SGD [1]}
9:     else
10:       $g_i \leftarrow \text{clip}(g_i, C_{\text{dp}})$  {same clip norm, no noise – isolates the privacy cost}
11:    end if
12:     $(\theta_i, \phi_i) \leftarrow (\theta_i, \phi_i) - \eta g_i$ 
13:  end for
14:   $\bar{\theta}, \bar{\phi} \leftarrow \frac{1}{N} \sum_{i=1}^N (\theta_i, \phi_i)$  {FedAvg [13], equal client weights}
15:  for each client  $i = 1, \dots, N$  do
16:     $(\theta_i, \phi_i) \leftarrow (\bar{\theta}, \bar{\phi})$  {broadcast}
17:  end for
18: end for
```

4.3 Experimental Setup

Both the centralised and federated campaigns use 300 training and 50 evaluation episodes per seed. The centralised campaign uses a fixed 30-seed list (seeds 900–929) with three arms: GAT-CTDE; an independent-DQN ablation (one per-gNB DQN instance, no parameter sharing, no topology information) that isolates the GAT/centralised- training contribution; and a single-agent DQN baseline run against the flattened joint state. All three arms use matched hyperparameters (discount factor $\gamma = 0.95$, per-episode epsilon decay). The federated campaign uses the first 10 of these 30 seeds and a five-level noise-multiplier sweep $\sigma \in \{0, 0.5, 1, 2, 4\}$, so its $\sigma=0$ level is directly comparable to the centralised campaign's same 10 seeds. All confidence intervals are 95% bootstrap percentile intervals (10,000 resamples); paired comparisons use the Wilcoxon signed-rank test.

5 Evaluation Safeguards

This section diagnoses a training collapse that a standard SLA-compliance metric could not detect on its own. We rely throughout on two correctness-aware metrics computed from data every run already logs: mean reward per step, the actual training objective, and block precision, the fraction of blocked requests that correctly target the slice the reward calibration makes reject-optimal. Section 5 below states why these are

necessary in full; here they let us tell a genuine training failure apart from a legitimate one.

While generating a privacy-utility sweep for the federated variant (Section 6.2), noise levels that should have produced genuinely different, stochastically-trained checkpoints instead produced bit-identical per-seed evaluation compliance. Tracing this back, we found that our first GAT-CTDE encoder’s eval policy blocked zero requests across every evaluation episode, on all 30 campaign seeds. An always-accept policy is a functional no-op, so any two checkpoints that reach it produce identical trajectories regardless of their internal weights. We confirmed, rather than assumed, that this was a genuine failure and not a legitimate optimum. A direct Q-value probe on the collapsed checkpoint showed $Q(\text{accept}) - Q(\text{reject})$ positive and large for every slice, including mMTC, while the independent-DQN ablation (Section 6.1), which sees identical raw features through no encoder at all, does not exhibit this collapse. This isolates the encoder, not the reward or training loop, as the fault.

We compared the encoder’s own embedding for a synthetic idle state against a congested one on the collapsed checkpoint. The raw input difference has ℓ_2 norm 2.55 and the resulting embedding difference has norm 19.6, an apparently large, congestion-sensitive response, but the cosine similarity between the two embeddings is 0.99997. This means the encoder was encoding congestion almost entirely as embedding *magnitude*, with no mechanism preventing this and no rotation of direction. Since the downstream Q-head is roughly linear in its input, larger magnitude pushes $Q(\text{accept})$ up faster than $Q(\text{reject})$ as congestion increases, which runs backwards from the reward’s intent.

We add a per-layer **LayerNorm** to the encoder, constraining every node’s embedding to zero-mean, unit-variance regardless of input magnitude. This is a real, verified improvement (collapse drops from 30/30 to 27/30 seeds) but not a complete fix. We then apply the per-slice Q-heads of Section 4.1, which reduces collapse further, to 9/30 seeds (Fig. 2); a further variant disabling LayerNorm’s learnable affine transform made results *worse* (0/3 pilot seeds differentiated versus 1/3 with it enabled) and was reverted. Not every plausible fix helped, and we report the one that did not.

A separate, unrelated bug was caught while building an evaluation-time disruption harness. Our logging utility opens its output file in append mode and never truncates it, and GAT-CTDE was retrained in place three times while diagnosing the collapse above, each time writing evaluation results to the same per-seed path. Every one of the 30 seeds’ evaluation logs had silently accumulated all three runs’ episodes stacked together, confirmed directly by each log’s own episode index resetting to 1 twice. Every GAT-CTDE number reported through an earlier draft of this paper, and every M3 result reusing GAT-CTDE’s centralised evaluation as a reference point, was therefore a hidden three-way average across two architecturally-obsolete runs and the true final one. The saved model checkpoints themselves were unaffected (a different save routine overwrites rather than appends, verified by reproducing one seed’s final-block statistics exactly from its checkpoint in isolation), so the fix required no retraining; we cleared every stale log, re-ran evaluation once against the already-trained checkpoints, and hardened the training script to clear its output path before any future run. All numbers in this paper reflect that corrected data.




Fig. 2 Fraction of the 30 campaign seeds reaching genuine, correctly-targeted differentiated shedding (mMTC-only blocking under congestion), across the original unnormalised encoder and the two fixes described in Section 5.

The fixes above make the architecture more correct while making its `sla_compliance_all_slices` score *worse* (Section 6.1), because blocking mMTC under congestion, the reward-optimal action, does not recover mMTC’s own SLA margin in this stress regime, so every newly-correct decision costs compliance the metric never credits back. We therefore report two further metrics, both computed from data every arm already logs: **mean reward per step**, the actual training objective; and **block precision**, the fraction of blocked requests that target mMTC specifically, the only slice the reward calibration ever makes reject-optimal (undefined, and reported as such, for seeds that never block anything). Any admission-control evaluation where correct behaviour cannot recover an already-stressed slice’s compliance should report a decision-level correctness metric alongside, not instead of, an outcome-level compliance one. This is the general lesson this specific collapse makes concrete.

This disagreement is not unique to the one collapse diagnosed above. Re-checking every paired comparison already reported in this paper’s centralised (Section 6.1), federated (Section 6.2), disruption (Section 6.3), and cluster-scaling (Section 7) campaigns against both metrics together, the correctness-aware pair and `sla_compliance_all_slices` disagree on direction or significance in 44 of 49 comparisons (89.8%) on the seed ranges already committed at the time, independently confirmed at 21/24 (87.5%) on a disjoint supplementary seed range drawn specifically




Fig. 3 M2 centralised campaign, per-seed evidence for the correctness-aware metrics (Section 5). (a) Paired per-seed comparison against single-agent DQN on mean reward per step. Each line is one seed, coloured by whether GAT-CTDE scored higher (green, 20), tied (grey, 3), or lower (red, 7), the source of this figure’s borderline $p = 0.0577$. (b) Full per-seed reward distribution for all three arms; GAT-CTDE is split into seeds that learned genuine differentiated shedding (blue, 21) versus seeds still stuck at the always-accept collapse (grey $\times$, 9), shown without claiming the split predicts reward, since several still-collapsed seeds score among the arm’s highest.

to check whether the first count was itself a one-off. `sla_compliance_all_slices` is not an occasionally-misleading proxy for this problem; on this evidence it is close to uninformative about decision quality on its own, and every number in this paper that matters is reported through the correctness-aware pair, not the compliance metric, for exactly this reason.

The second safeguard this paper applies throughout is a reproduction-versus-replication discipline: every headline paired comparison is recomputed on an independent, disjoint seed sample drawn after the first, and both figures are reported together rather than only the first. This is not a formality; applied to the disruption campaign (Section 6.3), it directly contradicted an initial reading that independent DQN was architecturally immune to agent churn, a claim that looked confirmed on the committed sample alone and was retracted once the independent sample showed a significant, larger-than-coordinated cost in the opposite direction. Section 6.3 reports the full tally of every headline comparison checked this way; the churn retraction is the one exception among them, not the rule.

6 Multi-gNB Results: Centralised, Federated, and Disruption

6.1 Centralised Campaign

The $N = 3$ fully-connected topology used throughout this campaign is deliberately the setting where a shared graph-attention encoder has the least topology to exploit,

Table 1 M2 campaign results, 30 seeds/arm, 95% bootstrap CIs.

Arm	Mean reward/step	Block precision	Compliance*
GAT-CTDE	14.062 [13.91, 14.23]	0.952	0.151 [0.083, 0.225]
Indep. DQN	12.601 [11.30, 13.49]	0.901	0.049 [0.013, 0.093]
Single-agent	13.867 [13.69, 14.05]	0.987	0.115 [0.061, 0.176]

* Low compliance is the metric artifact Section 5 explains, not a performance failure.

since adjacency across three gNBs is close to trivial at full connectivity. This makes it a clean substrate for the evaluation-integrity findings that follow, uncomplicated by any topology effect; the encoder’s coordination advantage is put to a genuine test only once the cluster scales to $N = 19$ under sparser adjacency in Section 7.

Table 1 and Fig. 3 report the final, post-fix 30-seed campaign. On `sla_compliance_all_slices`, GAT-CTDE’s mean is 0.151, below its own pre-fix number (0.353) since the metric penalises newly-correct mMTC-blocking decisions, and the paired comparison against single-agent DQN shows no clear direction, at +0.036 (95% CI $[-0.050, 0.126]$, Wilcoxon $p = 0.4549$, 11 wins / 9 ties / 10 losses). On mean reward per step, the metric that reflects the actual objective, GAT-CTDE beats the independent-DQN ablation decisively, by +1.461 (95% CI $[0.550, 2.782]$, $p < 0.0001$, 27/0/3), isolating a real contribution from the shared GAT encoder and centralised training, but its edge over single-agent DQN is small and not statistically decisive at this sample size, at +0.195 (95% CI $[-0.017, 0.442]$, $p = 0.0577$, 20/3/7). Its block precision (0.952) is high regardless, since when GAT-CTDE blocks a request, it is almost always the reward-correct slice. 21 of the 30 seeds show genuine, correctly-targeted differentiated shedding, up from 3/30 under the LayerNorm-only fix and 0/30 originally. This is a real, substantial behavioural improvement that does not, at this sample size, translate into a decisive reward edge over the single-agent baseline specifically.

6.2 Federated Training and Differential Privacy

Comparing the centralised campaign against the federated variant with no DP noise, on the same 10 seeds, mean reward per step moves from 14.073 to 13.940, a change of +0.133 (95% CI $[-0.138, 0.470]$, Wilcoxon $p = 0.8125$). This is no measurable cost to federating at this sample size, and both arms’ block precision is perfect (1.000) on seeds that block anything at all. Sweeping the DP noise multiplier (Fig. 4(b)), block precision is perfect through $\sigma = 1.0$ then drops sharply, to 0.834 at $\sigma = 2.0$ and 0.584 at $\sigma = 4.0$. This is a real threshold, not a smooth degradation, while mean reward per step stays in a comparatively narrow band across the whole sweep (12.45–14.19), since DP noise erodes *which* slice gets blocked, not whether the policy tries to differentiate at all. The compliance-based version of this same sweep (not shown) is non-monotonic with sign flips between adjacent σ levels. This is a second instance, independent of Section 5’s architectural finding, of `sla_compliance_all_slices` failing to track the quantity we actually care about. Read together, federation and the FedProx null result (Section 4) answer a deployability question directly: periodic, infrequent synchronisation is sufficient for this architecture, so a disaggregated RAN Intelligent Controller

Table 2 Training hyperparameters, all fixed across seeds.

Hyperparameter	Value
Learning rate (GAT-CTDE)	1×10^{-4}
Learning rate (DQN baselines)	1×10^{-3}
Discount γ	0.95
Epsilon: start $\rightarrow$ end (decay/episode)	$1.0 \rightarrow 0.05$ (0.985)
Target-network sync	every 10 episodes
Optimiser	Adam
Training / evaluation episodes	300 / 50
<i>GAT encoder (GAT-CTDE only)</i>	
Hidden / output dim	16 / 16
Attention heads	4
Encoder layers	2
Q-head hidden dim	32
<i>Federated + DP (federated GAT-CTDE only)</i>	
Local steps per round	50
Aggregator	FedAvg ($\mu=0$)
DP gradient clip norm	1.0
DP noise multiplier sweep σ	$\{0, 0.5, 1, 2, 4\}$

does not need per-step coordination across gNBs to match centralised training’s own reward, and differential privacy can be added at a bounded, known precision cost rather than an unpredictable one.

6.3 Disruption Resilience

We evaluate every arm’s already-trained, frozen checkpoint (no retraining) against three mid-episode disruptions injected at evaluation time: *gNB dropout* (one gNB’s requests are forced to reject and its row in the joint feature matrix is zeroed, for a window); *demand spike* (a transient arrival-rate multiplier for a fixed window); and *agent churn* (one gNB’s decisions are made by a fresh, never-trained copy of its own policy class instead of the loaded checkpoint, for a window). Agent churn is a different, well-motivated failure mode from dropout (active-but-uncoordinated vs. absent) that applies to every genuinely multi-agent arm; single-agent DQN is excluded. Severity is the window’s duration (10%/30%/60% of the episode) for dropout and churn, and the arrival-rate multiplier ($2\times/4\times/8\times$) for spike. We normalise reward by each step’s own pending-request count before averaging, since raw reward mechanically inflates with request volume regardless of decision quality.

Each disruption kind intercepts the frozen policy’s per-step evaluation loop differently: dropout acts twice, corrupting the target gNB’s input row and separately forcing that gNB’s own requests to reject, so a policy cannot simply learn around the corrupted feature; agent churn substitutes a fresh, never-trained policy’s actions for the target gNB only, leaving every other agent’s decisions untouched; demand spike is entirely external to the policy, acting on the arrival process the environment synthesises for the following step. All three are evaluation-only perturbations, and the policy itself is never updated.




Fig. 4 M3 federated and differential-privacy sweep, per-seed evidence. (a) Paired per-seed federation cost on mean reward per step, same 10 seeds. Centralised scores higher on 2 seeds, federated on 3, 5 ties, no measurable difference (Wilcoxon $p = 0.8125$). (b) Privacy-utility curve using block precision against the DP noise multiplier σ , individual seed values (light dots) behind the mean, with the seed count that ever blocked anything at each σ printed below (precision is undefined otherwise). Precision is perfect through $\sigma = 1.0$, dropping sharply by $\sigma = 4.0$.




Fig. 5 M4 disruption campaign, 10 seeds/arm. (a)-(b) Normalised reward cost (baseline - disrupted) vs. severity for dropout and agent churn. Every arm, every severity, shows a real cost that accelerates with severity rather than growing linearly. (c) Block precision vs. demand-spike multiplier: flat and near-ceiling for every arm, unlike (a)-(b), since spike does not threshold on the metric that matters.

Dropout costs every arm significantly at every severity (Wilcoxon $p \leq 0.03$, 9 or 10 of 10 seeds hurt in every case), and the cost accelerates with severity rather than growing proportionally. For GAT-CTDE, the cost is $+0.165$ at the 10% window, $+0.601$ at 30%, and $+1.710$ at 60%, and the same super-linear pattern holds for every arm. Agent churn shows the same accelerating shape for GAT-CTDE ($+0.059/+0.200/+0.529$, $p = 0.0039$ throughout) and the federated arm ($+0.078/+0.246/+0.585$, $p = 0.0039$

throughout; regenerated post-hoc from the still-intact frozen checkpoint after the original per-seed logs were found missing during this paper’s own reproducibility audit, with environment drift over the intervening months identified as the cause by systematic elimination of every code- and data-level alternative). Independent DQN’s own churn cost needs both samples reported together to state honestly. On the seeds used throughout the rest of this paper (900–909), it is small and not conventionally significant (+0.000/+0.168/+0.640, $p = 0.084/0.106/0.065$), which an earlier pass read as independent DQN being architecturally immune to churn, since it never attends to another agent’s features. An independent-seed replication (seeds 1000–1009, disjoint from every other result in this paper) contradicts that reading directly. Independent DQN’s churn cost there is significant at every severity (+0.097/+0.657/+1.223, $p = 0.0020$ throughout) and *larger* than GAT-CTDE’s or the federated arm’s at matched severities, the opposite of immune. We do not conclude independent DQN is immune to churn, since the direction was consistent across both samples (churn always hurts, never helps). What was wrong was treating a borderline, underpowered $n = 10$ result as a confirmed architectural property, and this section reports the corrected reading, not the retracted one, as this paper’s discipline of checking every headline result against an independent seed sample, not the same seeds reproduced again, requires.

Demand spike shows no threshold at all on block precision, which stays flat and near-ceiling through the highest tested multiplier for every arm; a reward-based signal that spike *improves* per-step reward is real and reproduces across both seed samples for GAT-CTDE, independent DQN, and the federated arm ($p \leq 0.014$ in both), but traces to a fixed per-step violation penalty diluted by a larger request-volume denominator, not better decision-making. Block precision, a volume-invariant ratio, is the metric that actually answers whether spike corrupts decision quality, and it says no. Single-agent DQN’s version of this same signal is significant in the original sample ($p = 0.0020$) but not the replication ($p = 0.084$); the direction is consistent, the significance is sample-dependent, and both are reported as such.

Every dropout and churn result above was independently checked against the 1000–1009 sample and holds up in direction and approximate magnitude. For example, GAT-CTDE’s dropout costs there are +0.170/+0.552/+1.449, close to the 900–909 sample’s +0.165/+0.601/+1.710. Every headline paired comparison in this paper was checked this way, recomputed from raw logs on both samples: GAT-CTDE’s edge over independent DQN (Section 6.1) holds decisively on the replication sample too (+0.767, 95% CI [0.391, 1.233], $p < 0.0001$, versus the committed sample’s +1.461 [0.550, 2.782]); its edge over single-agent DQN stays non-decisive on both (+0.445 [−0.049, 1.280], $p = 0.111$, versus +0.195 [−0.017, 0.442], $p = 0.058$); federation’s reward cost stays statistically indistinguishable from zero on both, with the point estimate’s sign flipping between samples (−0.080 [−0.241, 0.001] on replication versus +0.133 [−0.138, 0.470] committed) but neither ever significant. Of the seven headline paired comparisons checked this way across the whole paper, six show overlapping confidence intervals across samples; the one exception, independent DQN’s churn cost, is the retraction above.

A formal paired Wilcoxon test on (GAT-CTDE’s cost – independent DQN’s cost), computed across the same 10 seeds for every (kind, severity) cell, reaches significance only once, at dropout at the mildest severity (10% window), where GAT-CTDE’s cost is significantly *higher* than independent DQN’s (+0.0135, 95% CI [0.0053, 0.0236], Wilcoxon $p = 0.0059$, 9 losses / 0 ties / 1 win). Every other cell, dropout at 30%/60% and churn at all three severities, shows no significant difference ($p \geq 0.23$). We report the result as it is, real at one severity and absent everywhere else tested, not a general claim that coordination costs more under disruption.

7 Cluster-Size Scaling

We extend the cluster from $N = 3$ to $N \in \{7, 19\}$ under three adjacency structures (fully-connected, ring, hex-grid), scaling synthetic arrivals with N to hold offered load per gNB roughly constant. GAT-CTDE is the only arm whose Q-head consumes this adjacency; independent DQN and single-agent DQN never read a graph structure and serve as a topology-invariance check. We normalise reward per gNB before averaging (the same fix Section 6.3 applies on the severity axis) and report block precision alongside, since `sla_compliance_all_slices` is OR-aggregated across gNBs and trends toward 0 as N grows regardless of policy quality. This is a third instance of this paper’s central evaluation-integrity finding, now on the cluster-size axis. The primary sample is 12 seeds (900–911) at $N = 19$, all three topologies, all three arms; an independent-seed replication (1000–1002, disjoint) checks generalisation.

At the full 12-seed sample, the paired per-gNB reward-margin comparison at $N = 19$ is not significant for any topology (fully-connected -0.175 [95% CI $-0.528, +0.076$], $p = 0.7422$; ring -0.704 [$-1.879, +0.042$], $p = 0.7422$; hex -0.695 [$-1.873, +0.053$], $p = 0.9102$). This is mechanically explicable, not a red flag, since single-agent DQN’s always-accept collapse (below) collects the reward’s service term on every request with no rejection cost, while GAT-CTDE pays a real cost every time it correctly rejects mMTC. This paper’s evidence for GAT-CTDE’s advantage at $N = 19$ lives in collapse resistance and block precision, not the reward-margin metric.

Single-agent DQN collapses totally at $N = 19$: 12/12 primary-sample cells and 3/3 replication cells, every topology, zero exceptions (15/15 combined). This is confirmed at the individual-decision level by a direct Q-value probe identical in method to Section 5’s original diagnosis. GAT-CTDE collapses only partially, and independent DQN never fully collapses (0/36) but never cleanly differentiates either, sitting in a narrow, mediocre 0.29–0.49 block-precision band identical across all three topologies, as expected for an arm that never consumes adjacency (Section 9 tests this same claim again at $N=7$ under an independent replication sample and finds a small, rare exception to “never”). GAT-CTDE’s own non-collapsed seeds show a topology-dependent secondary failure mode, since most hold near-perfect precision, but at least one seed per topology drops to 0.0–0.5, including one seed showing perfect precision (1.000) at fully-connected and near-zero (0.000) at ring/hex on the identical seed, environment stream, and initial weights.

GAT-CTDE’s own collapse rate initially disagreed sharply between the primary sample (31%) and the replication sample (78%); rather than report the range as open,




Fig. 6 $N = 19$ cluster, per-seed evidence across independent seed samples. (a) Collapse rate (fraction of eval runs with zero blocks), pooled across all three topologies. Single-agent DQN collapses totally in every sample; GAT-CTDE’s rate varies sample-to-sample (per-sample breakdown in the body text), converging to a combined, seed-level-bootstrapped estimate of 35.4% [25.5%, 45.8%] across all 64 independent seeds (diamond marker, error bar). (b) Block precision by topology, primary sample, non-collapsed seeds only (single-agent DQN has none to show). GAT-CTDE’s non-collapsed seeds are mostly near-ceiling but with a low-precision seed at every topology, while independent DQN sits in a narrow, mediocre band, identical across topologies as expected for an arm that never consumes adjacency.

we drew a third, independent sample targeted specifically at this question (49 seeds, disjoint from both prior samples). Pooling all three samples gives 68/192 cells collapsed, 35.4% (seed-level bootstrap 95% CI [25.5%, 45.8%] across the 64 independent seeds, since collapse status correlates within a seed across its three topologies). GAT-CTDE’s collapse rate at $N = 19$ is best understood as roughly 35% with substantial seed-to-seed variability, and none of the three samples was wrong, since each was an honest reading of a genuinely wide distribution, and no single small sample should be trusted alone.

8 Live Single-gNB Deployment

Every multi-gNB, federated, and disruption result above is offline; this section asks what a real OAI testbed can and cannot establish about the offline-trained policy, and treats the answer to *cannot* as seriously as the answer to *can*. This rig cannot yet sustain a live multi-gNB test, so GAT-CTDE’s coordination benefit and the federated/DP results remain architecturally untestable on it; what follows uses a single-agent-DQN checkpoint (seed 900, dimension-matched to the live evaluation path) against a real gNB, real UEs, and real per-slice UDP traffic, at native load (1 UE/slice) and at double native load (2 UEs/slice).

An early live campaign on this rig read a series of findings from data collected while a scheduler-configuration bug was active: `sac1b.live.yaml`’s per-slice PRB-ratio

ranges, between commits 82b20f6 (2026-07-16) and 40ae2f3 (2026-08-19), converted to raw PRB grants below the gNB scheduler’s hard 5-PRB minimum grant size on all three slices, a cross-layer interaction between a configuration file and the scheduler’s own grant floor that a config-level review alone could not have caught. This produced catastrophic RLC retransmission failures indistinguishable, from the policy’s own vantage point, from genuine per-slice capacity scarcity. Fixed at its root (commit cda9d65, 2026-09-05, rescaling the config’s ratio ranges to clear the floor) rather than worked around, every result below was collected after the fix, with the scheduler’s own state directly instrumented and confirmed clear of the floor throughout every run reported here (zero below-floor grants across more than 100,000 combined scheduler-state readings across all live runs in this section).

The re-collected picture is a genuine deployability result, not an admission-control result. A dedicated live capacity probe, native traffic and a deliberately wide-open admission ceiling so the reading reflects the environment rather than a policy, found all three slices serve essentially 100% of offered demand with 0.000% loss, at both 1 and 2 UEs per slice (URLLC 100%, eMBB ≈ 99.9 –100%, mMTC ≈ 105 –107% of its own bursty-average demand, i.e. comfortably served). This rig, at every native load level it has been tested at, has never been shown to reach a point where any slice’s real admission decision carries a genuine capacity consequence. Taking the frozen checkpoint live at native load, its offline-trained decision logic executes correctly end to end: real per-slice traffic flows through the real E2 control loop, real accept/reject decisions are issued every step, and the policy’s characteristic low-priority-only blocking pattern (mMTC targeted, URLLC/eMBB untouched) reproduces at both 1 and 2 UEs per slice, with no catastrophic failure at either load. But because the environment itself never reaches scarcity, this blocking pattern cannot be read as contention management succeeding; it is at least as consistent with the policy exercising its own trained priority bias (mMTC carries the lowest `priority_weight` by the reward’s own calibration, Section 6.1) in a load regime that never actually forces the choice. **The honest claim this rig supports is deployability: the trained policy runs correctly, end to end, on real hardware. Whether its admission decisions are actually managing real contention is a question this rig’s own achievable load cannot yet put to the test**, and is instead answered by the disclosed offline stress environment used throughout Sections 6.1–6.3, 7, and 9, deliberately calibrated to a harder congestion range than this rig has ever been measured to reach.

A separate, standing finding survives this re-framing unaffected: a genuine unit-mismatch code defect in the live SLA-margin diagnostic, where the margin computation divides a per-slice backlog reading by a calibration constant tuned for the offline environment’s synthetic proxy, while live the identically-named field is wired to the real MAC scheduler’s transmit-buffer byte count, two to three orders of magnitude larger in scale. Re-collected at native load post-fix, this no longer manifests as a catastrophic reading (eMBB’s margin now lands at -47.64 , roughly $21,000\times$ smaller in magnitude than an earlier catastrophic $-1,002,377.5$ reading, since the backlog level that had driven that earlier number was itself very likely the pre-fix floor bug’s own starvation signature, not organic contention) but the unit mismatch in how the diagnostic is computed is real regardless of load, and the fix already applied, converting

backlog to a physically meaningful queueing-delay estimate via Little’s Law against each slice’s own latency budget rather than an arbitrary rescaling constant, remains correct and is kept.

An earlier version of this section reported a targeted simulator recalibration (anchoring offline served capacity to two hardcoded, UE-count-doubling live measurements) that appeared to reverse a complete live collapse at double load. Re-collection under the fixed scheduler config shows that collapse did not reproduce at all once the floor bug was fixed at its root; the recalibration’s own anchors described the floor bug’s signature, not a real physical capacity constraint, and it is retired rather than carried forward (the anchors themselves are preserved, unmodified, as evidence of the misdiagnosis, not deleted). A companion submission [10] separately measures the E2 control loop itself as negligible overhead (median round-trip 0.57 ms) from a policy under 20,000 parameters, on the same rig – a live feasibility result this deployability finding is consistent with and does not depend on.

9 Cluster-Size Scaling at $N=7$

§7’s own $N=7$ number was a 3-seed pilot, never resampled at scale because that work judged $N=19$ ’s collapse question the higher-value target: GAT-CTDE and single-agent DQN both held 1.000 precision in that pilot (0/3 collapsed each). We closed this gap with a properly powered resample under the identical, unmodified stress environment used throughout §7 (`ClosedLoopKpmSource`, no recalibration of any kind), 12 seeds primary (900–911, matching §7’s own primary identity) plus an independent 12-seed replication (1000–1011, disjoint), three topologies, three arms, the same 100-train/20-eval budget the original pilot used.

Table 3 reports both samples directly. The pilot’s finding does not survive: single-agent DQN’s clean $N=7$ collapse rate is 91.7% [75.0%, 100%] primary and 75.0% [50.0%, 100%] replication, both far above the pilot’s 0%, and this is the single clearest result in this scaling extension, not merely a wider interval around the pilot’s point estimate but the opposite conclusion from a properly powered sample. GAT-CTDE’s $N=7$ rate (41.7% [16.7%, 66.7%] primary, 58.3% [33.3%, 83.3%] replication) is statistically indistinguishable from its own $N=19$ pooled estimate (35.4% [25.5%, 45.8%]), meaning its collapse-proneness is present from $N=7$ onward rather than an effect that only emerges at large cluster size. Independent DQN’s primary sample matches the $N=19$ pooled reading exactly (0/12, never collapses); its replication sample does not (2/12, 16.7% [0.0%, 41.7%]) – the same small, real, non-zero exception this paper’s $N=19$ result already shows a hint of, now observed independently at a second scale. Independent DQN’s own “never fully collapses” claim is best read as “rarely, not never” at both scales, not retracted at either.

This closes the one gap §7 itself flagged (an under-powered pilot, not a properly-sampled result) using the same environment throughout, with no recalibration of any kind: the original, unmodified stress regime already used at $N=19$.

Table 3 $N=7$ collapse rate, clean stress environment (`ClosedLoopKpmSource`, unmodified), 12 seeds/sample, 3 topologies $\times$ 3 arms, 95% bootstrap CI.

Arm	Primary (900–911)	Replication (1000–1011)	Original 3-seed pilot
Single-agent DQN	91.7% [75.0, 100]	75.0% [50.0, 100]	0% (0/3)
GAT-CTDE	41.7% [16.7, 66.7]	58.3% [33.3, 83.3]	0% (0/3)
Independent DQN	0.0% [0.0, 0.0]	16.7% [0.0, 41.7]	0% (0/3)

10 Conclusion

We set out to test whether a standard evaluation metric, and a live testbed’s own evidentiary reach, can be trusted in multi-agent, privacy-constrained, disruption-prone slice admission control, and found neither can go unchecked. A standard SLA-compliance metric, an unguarded append-only log, an unseeded weight-initialisation path, and a cross-layer scheduler-configuration bug together hid four distinct failures: a training collapse the compliance metric scored no worse than correct behaviour, a hidden eval-data contamination bug, a same-seed-only reproducibility gap, and a live scheduler-floor misconfiguration that had silently manufactured the appearance of genuine contention in earlier live data. None of the four would have been caught by a standard pipeline without this paper’s correctness-aware metrics, reproduction-replication discipline, and direct live instrumentation of the scheduler’s own state.

GAT-CTDE, the multi-agent architecture that served as this evaluation’s substrate, delivers a decisive, statistically significant improvement over an independent-learner ablation on the correctness-aware reward metric, though its edge over our existing single-agent baseline specifically is small and not decisive at this sample size. This is an honest, non-inflated architectural finding, not the headline result this paper argues evaluation practice should not be built around alone.

Extending to federated training with differential privacy, federation adds no measurable reward cost, and privacy is costly only past a real threshold in decision precision, not a smooth trade-off; FedProx earns no measurable dividend under genuine per-gNB load heterogeneity at any tested strength, traced to real per-round client drift being too small relative to the loss scale to matter, even at ten times the round length. Under mid-episode gNB dropout, demand spikes, and agent churn, dropout and churn produce a real, accelerating cost with severity, while demand spikes do not threshold at all on the metric that matters. Cross-checking every result against an independent seed sample, not just the same seeds reproduced again, caught one architectural claim (that independent DQN is immune to agent churn) that a single sample had made look confirmed but was not, and it is retracted here rather than reported.

Scaling the cluster to 19 gNBs, single-agent DQN’s collapse becomes total and GAT-CTDE’s remains only partial, a real cluster-size-specific effect; GAT-CTDE’s own collapse rate, read at two disagreeing values across two small samples (31% vs. 78%), was resolved rather than left open with a third, independently-drawn sample, converging on 35.4% [25.5%, 45.8%] across 64 seeds, a rate a properly powered $N=7$ resample under the identical environment shows is already present at the smaller

scale rather than an $N=19$ -specific phenomenon. Finally, a live single-gNB deployment establishes what this rig can and cannot show: the offline-trained decision logic executes correctly end to end on real hardware, a genuine deployability result, but a cross-layer scheduler-configuration bug had earlier manufactured the appearance of live contention this rig cannot actually produce; fixed at its root and re-collected, this rig has never been shown to reach a load at which any slice’s admission decision carries a real capacity consequence, so admission-control-under-contention remains a claim this paper makes only in the disclosed offline stress regime, not live.

Four limits bound every claim in this paper. First, every multi-gNB, federated, and disruption result is offline-only, since this rig cannot yet sustain a live multi-gNB test, so whether these rankings would survive contact with a live, multi-gNB testbed remains untested, not just under-tested. A companion submission [10] establishes that this rig’s live E2 control loop is itself well-behaved for a single gNB, negligible round-trip latency, a lightweight policy, but whether that same loop stays well-behaved once it is coordinating $N > 1$ gNBs live is exactly the untested step this limit names, not something either paper yet answers. Second, the offline stress environment used throughout is deliberately calibrated harder than any congestion level this live rig has been measured to reach, a disclosed design choice reported as such rather than validated against a reachable live regime; the admission-control-under-contention claims in this paper are made in that disclosed regime, not live. Third, the graph representation models only the gNB cluster; it does not include multi-access edge computing nodes or any compute-placement dimension, so any extension toward a fuller realisation of a graph that also represents edge compute remains future work. Fourth, this paper’s federated scenario is a synthetic privacy setting (one operator’s cluster, not genuine cross-organisation data separation), so we lead with the operational benefit of periodic, not per-step, synchronisation for disaggregated RAN Intelligent Controllers rather than a privacy motivation this scale cannot substantiate. Together, these results argue that admission-control evaluation in multi-agent, privacy-constrained, disruption-prone settings needs metrics tied to the actual training objective, not only to a downstream SLA proxy that can reward the wrong policy, that architectural claims need more than one seed sample before they are trusted, and that a live testbed’s own evidentiary reach needs to be checked, not assumed, before its silence about contention is read as safety. A companion paper in preparation [11] examines a wider set of live evaluation-integrity findings from this same rig in more depth; the results reported here stand independently of it.

Acknowledgements. This work was supported by the Fundamental Research Grant Scheme (FRGS) of the Ministry of Higher Education (MOHE), Malaysia, under grant FRGS/1/2024/ICT06/TAYLOR/02/1.

Declarations

The following statements should be understood to apply to this manuscript unless otherwise indicated.

- **Funding:** This work was supported by the Fundamental Research Grant Scheme (FRGS), Ministry of Higher Education (MOHE), Malaysia, under grant number FRGS/1/2024/ICT06/TAYLOR/02/1.
- **Competing interests:** The authors declare no competing interests.
- **Ethics approval and consent to participate:** Not applicable, since this study involves no human participants, animal subjects, or identifiable patient data; all reported results are either offline simulation campaigns or measurements from a laboratory radio testbed.
- **Consent for publication:** Not applicable, since this manuscript contains no personal data or identifying information about any individual.
- **Data availability:** The experiment logs, checkpoints, and campaign result files underlying this paper’s findings ($\approx 48.6\text{GB}$) are archived on Zenodo at [DOI to be minted on deposit] under a CC-BY-4.0 licence, split as an offline-campaign deposit (the centralised, federated, disruption, and cluster-scaling training and evaluation logs, $\approx 40\text{GB}$) and a live re-collection deposit (the single-gNB deployment logs, $\approx 8.5\text{GB}$). Every number reported in this paper was independently re-derived from this archived data against a pinned software environment (Python 3.12.3, PyTorch 2.13.0+cu130, NumPy 2.5.1, SciPy 1.18.0) as part of a dedicated reproducibility audit; we treat this pinning, and the audit itself, as a release strength rather than a formality, since it is what let us catch and either fix or honestly disclose every discrepancy between an originally reported number and what the archived data actually supports.
- **Materials availability:** Not applicable, since this study involves no physical or biological materials.
- **Code availability:** The framework and experiment scripts used to produce this paper’s results, including the reproducibility manifest and analysis scripts used to independently re-derive every table and figure in this paper from the archived data above, are available at [GitHub URL] under [licence to be chosen].
- **Author contribution:** M. P. Kunasegran: conceptualisation, methodology, software, investigation, writing (original draft). W. L. Pang: supervision, writing (review and editing). S. K. Phang: supervision, writing (review and editing).

References

- [1] Abadi M, Chu A, Goodfellow I, et al (2016) Deep learning with differential privacy. In: Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS), pp 308–318
- [2] Ahmadi M, Moayyedi A, Sulaiman M, et al (2024) Generalizable 5G RAN/MEC slicing and admission control for reliable network operation. IEEE Trans Netw Service Manag <https://doi.org/10.1109/TNSM.2024.3437217>
- [3] Bouroudi A, Outtagarts A, Hadjadj-Aoul Y (2025) GNN-based multi-agent DRL for energy-efficient multi-domain 6G resource orchestration. In: Proc. IEEE Global Commun. Conf. (GLOBECOM), pp 4457–4462, <https://doi.org/10.1109/GLOBECOM59602.2025.11432814>

- [4] Bun M, Steinke T (2016) Concentrated differential privacy: Simplifications, extensions, and lower bounds. In: Theory of Cryptography Conference (TCC), pp 635–658
- [5] Fatehi K, Rahmani Ghourtani M, Sonee A, et al (2026) Interpretable attention-based multi-agent PPO for latency spike resolution in 6G RAN slicing. In: Proc. IEEE Int. Conf. Commun. (ICC), <https://doi.org/10.1109/ICC59461.2026.11586913>
- [6] Gong Y, Gong P, Jiang H, et al (2026) Intelligent resource allocation for multiaccess edge computing in 5G ultra-dense slicing network using federated multiagent DDPG. China Commun 23(1):273–289. <https://doi.org/10.23919/JCC.fa.2022-0187.202601>
- [7] Kunasegran MP, Pang WL, Phang SK (2025) Optimising resource allocation in 5G networks: Balancing URLLC and eMBB traffic under gNB congestion. In: 2025 Multimedia University Engineering Conference (MECON), pp 1–6
- [8] Kunasegran MP, Pang WL, Phang SK (2025) SLA-aware DRL for joint slice admission control and load balancing in multi-service 5G RAN. In: 2025 9th International Conference on Recent Advances and Innovations in Engineering (ICRAIE), pp 207–212
- [9] Kunasegran MP, Pang WL, Phang SK (2026) Comprehensive review of optimisation techniques for user-centric distributed network slicing in 5G networks. IEEE Access 14:114630–114648. <https://doi.org/10.1109/ACCESS.2026.3716911>
- [10] Kunasegran MP, Pang WL, Phang SK (2026) DQN-QoE slice admission control for live OAI open RAN. In: 2026 International Automatic Control Conference (CACs), submitted
- [11] Kunasegran MP, Pang WL, Phang SK (2026) Evaluation integrity in live multi-agent radio access network control: A systems study. Companion work, in preparation
- [12] Li T, Sahu AK, Zaheer M, et al (2020) Federated optimization in heterogeneous networks. In: Proceedings of Machine Learning and Systems (MLSys), pp 429–450
- [13] McMahan B, Moore E, Ramage D, et al (2017) Communication-efficient learning of deep networks from decentralized data. In: Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS), pp 1273–1282
- [14] Mohajer A, Mirzaei A, Darabi M, et al (2026) Joint SLA-aware task offloading and adaptive service orchestration with graph-attentive multi-agent reinforcement learning. IEEE Trans Netw Service Manag 23:3737–3754. <https://doi.org/10.1109/TNSM.2026.3673188>

- [15] Qasim M, Imran M, Akram F, et al (2025) Transformer and graph neural network-based federated learning for cellular traffic prediction with sustainability analysis. *IEEE Access* 13:159649–159666. <https://doi.org/10.1109/ACCESS.2025.3605380>
- [16] Shao Y, Li R, Hu B, et al (2021) Graph attention network-based multi-agent reinforcement learning for slicing resource management in dense cellular network. *IEEE Trans Veh Technol* 70(10):10792–10803. <https://doi.org/10.1109/TVT.2021.3103416>
- [17] Sulaiman M, Moayyedi A, Ahmadi M, et al (2023) Coordinated slicing and admission control using multi-agent deep reinforcement learning. *IEEE Trans Netw Service Manag* 20(2):1110–1124. <https://doi.org/10.1109/TNSM.2022.3222589>
- [18] Tashman DH, Cherkaoui S (2026) Adversarial attacks in AI-driven RAN slicing: SLA violations and recovery. In: *Proc. Int. Wireless Commun. and Mobile Computing (IWCMC)*, <https://doi.org/10.1109/IWCMC69287.2026.11580033>
- [19] Veličković P, Cucurull G, Casanova A, et al (2018) Graph attention networks. In: *International Conference on Learning Representations (ICLR)*
- [20] Yasin TV, Yu CM, Wang LC (2025) Differential privacy federated edge learning-assisted for securing RAN intelligent controller in O-RAN 6G communications. In: *Proc. IEEE VTS Asia Pacific Wireless Commun. Symp. (APWCS)*, <https://doi.org/10.1109/APWCS67981.2025.11151868>
- [21] Zhang H, Zhou H, Erol-Kantarci M (2022) Federated deep reinforcement learning for resource allocation in O-RAN slicing. In: *Proc. IEEE Global Commun. Conf. (GLOBECOM)*, pp 958–963, <https://doi.org/10.1109/GLOBECOM48099.2022.10001658>