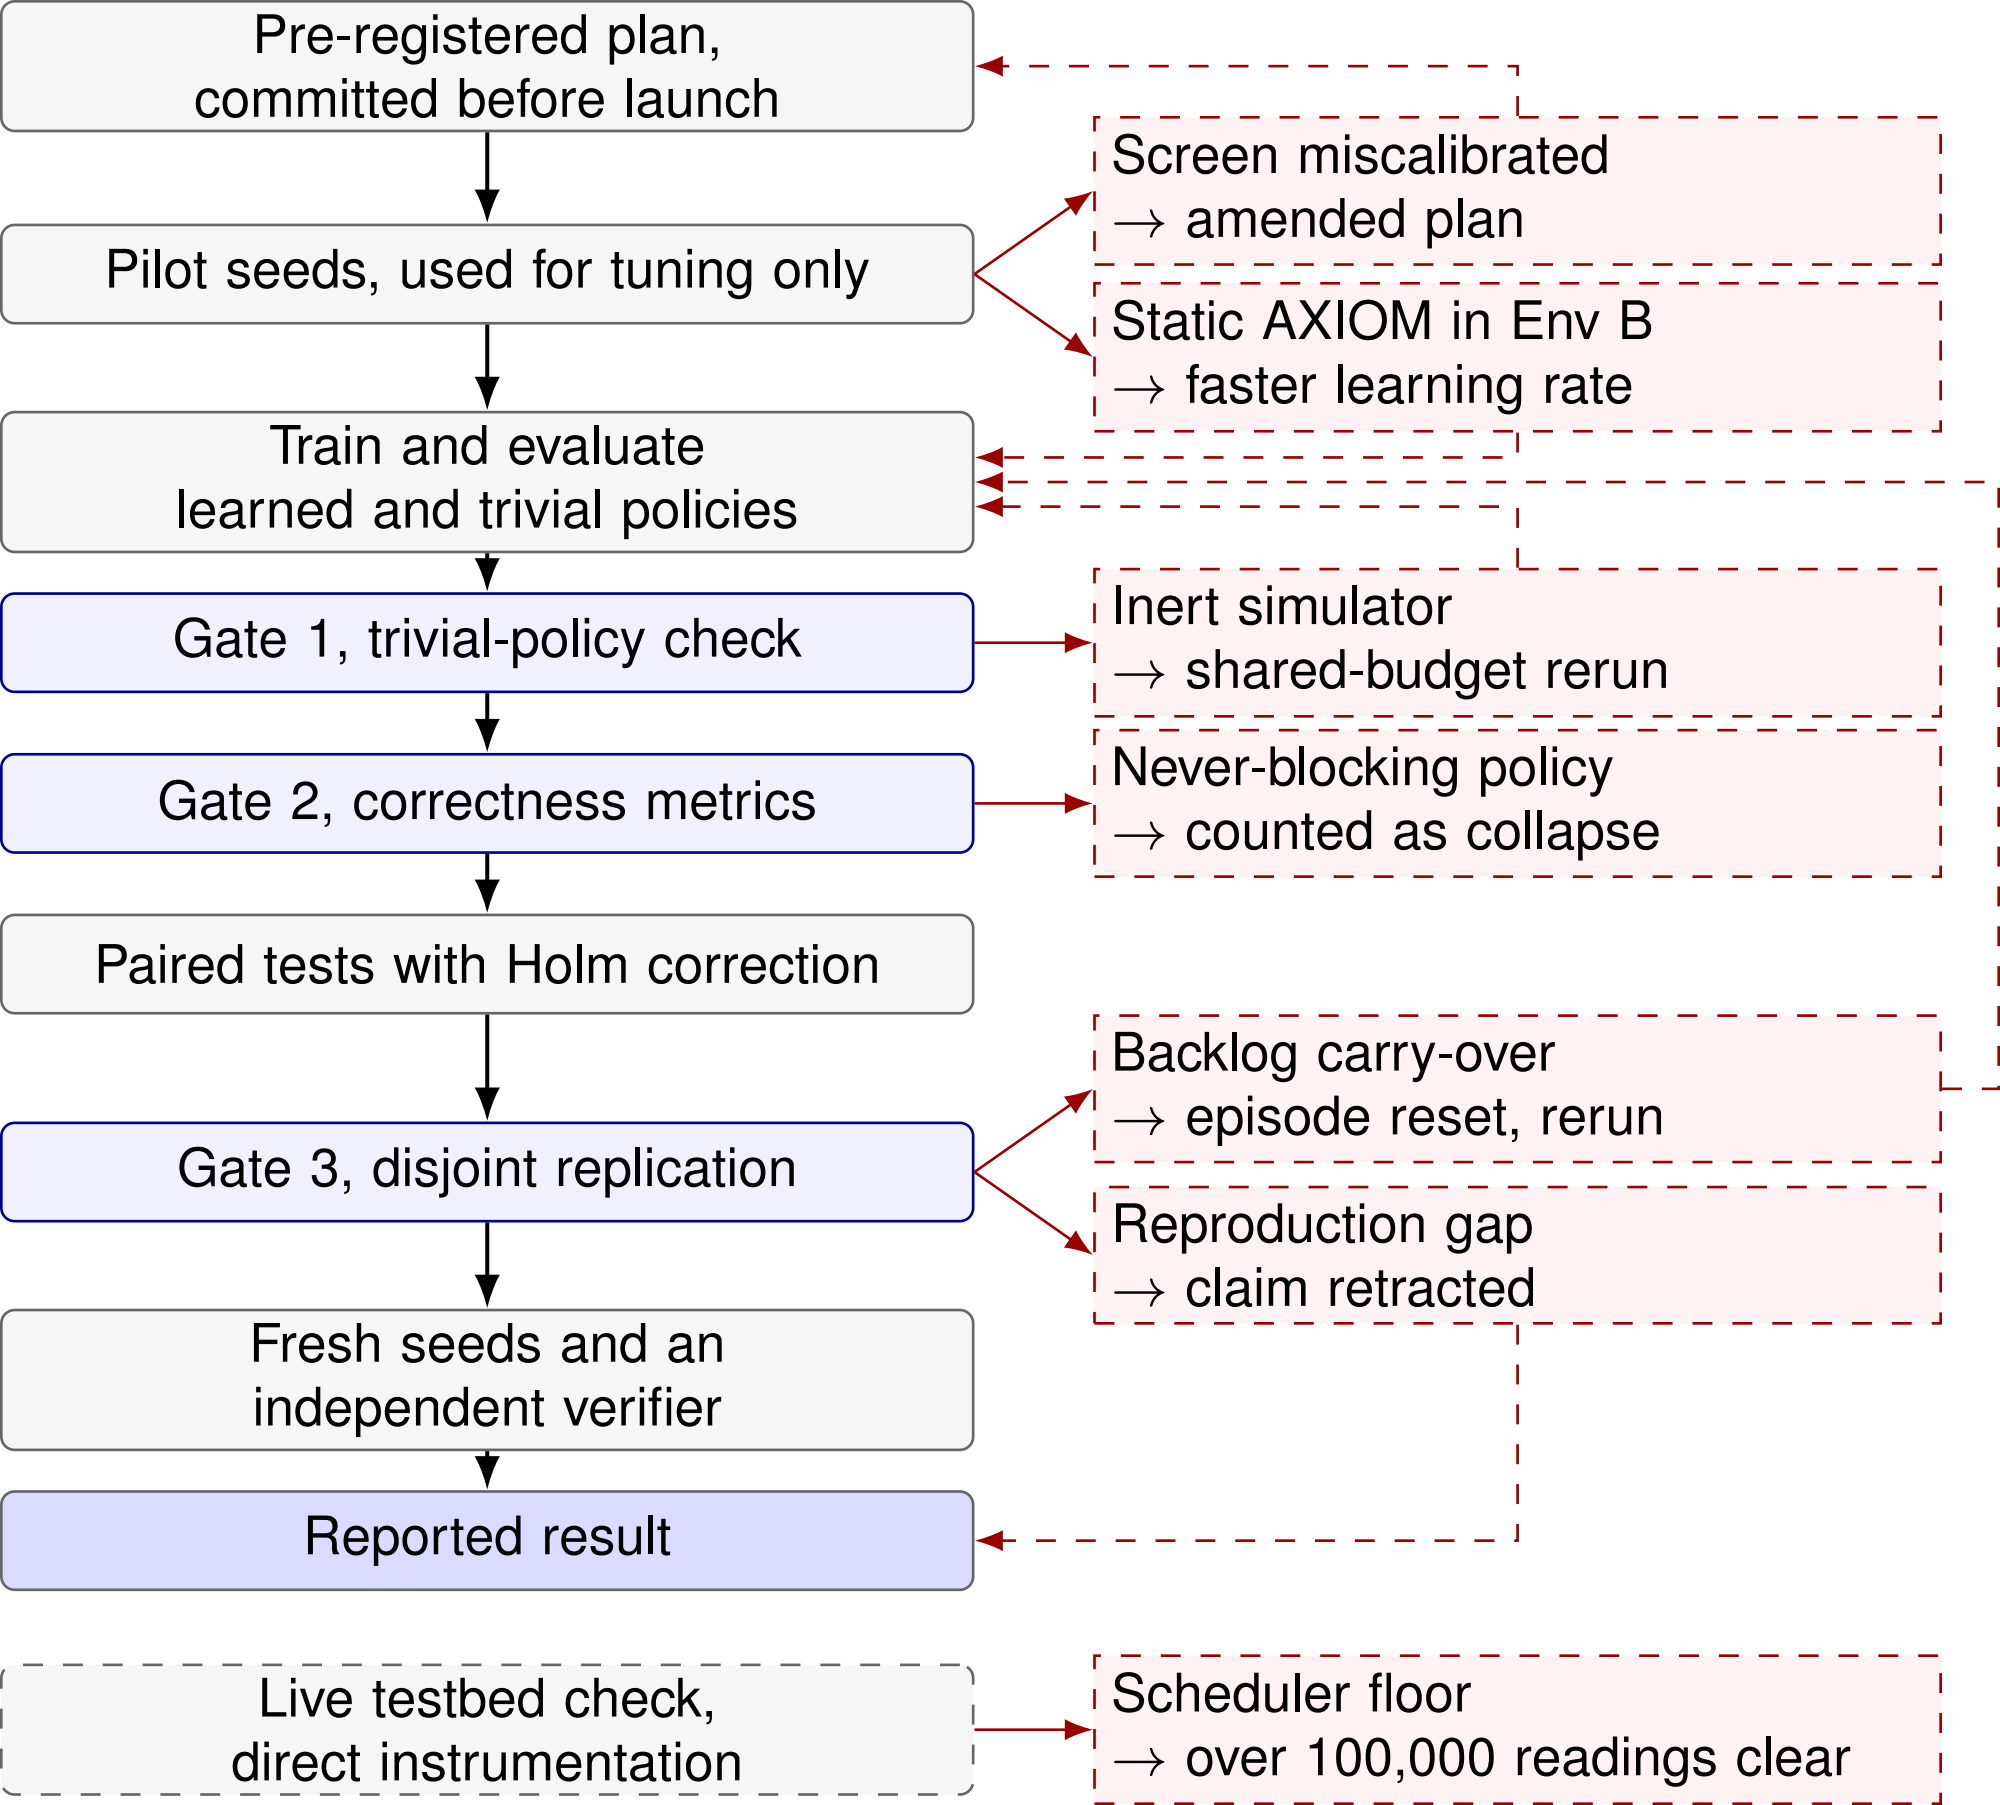
Pre-registered plan, committed before launch
Pilot seeds, used for tuning only
Train and evaluate learned and trivial policies
Gate 1, trivial-policy check
Gate 2, correctness metrics
Paired tests with Holm correction
Gate 3, disjoint replication
Fresh seeds and an independent verifier
Reported result
Live testbed check, direct instrumentation
Screen miscalibrated → amended plan
Static AXIOM in Env B → faster learning rate
Inert simulator → shared-budget rerun
Never-blocking policy → counted as collapse
Backlog carry-over → episode reset, rerun
Reproduction gap → claim retracted
Scheduler floor → over 100,000 readings clear