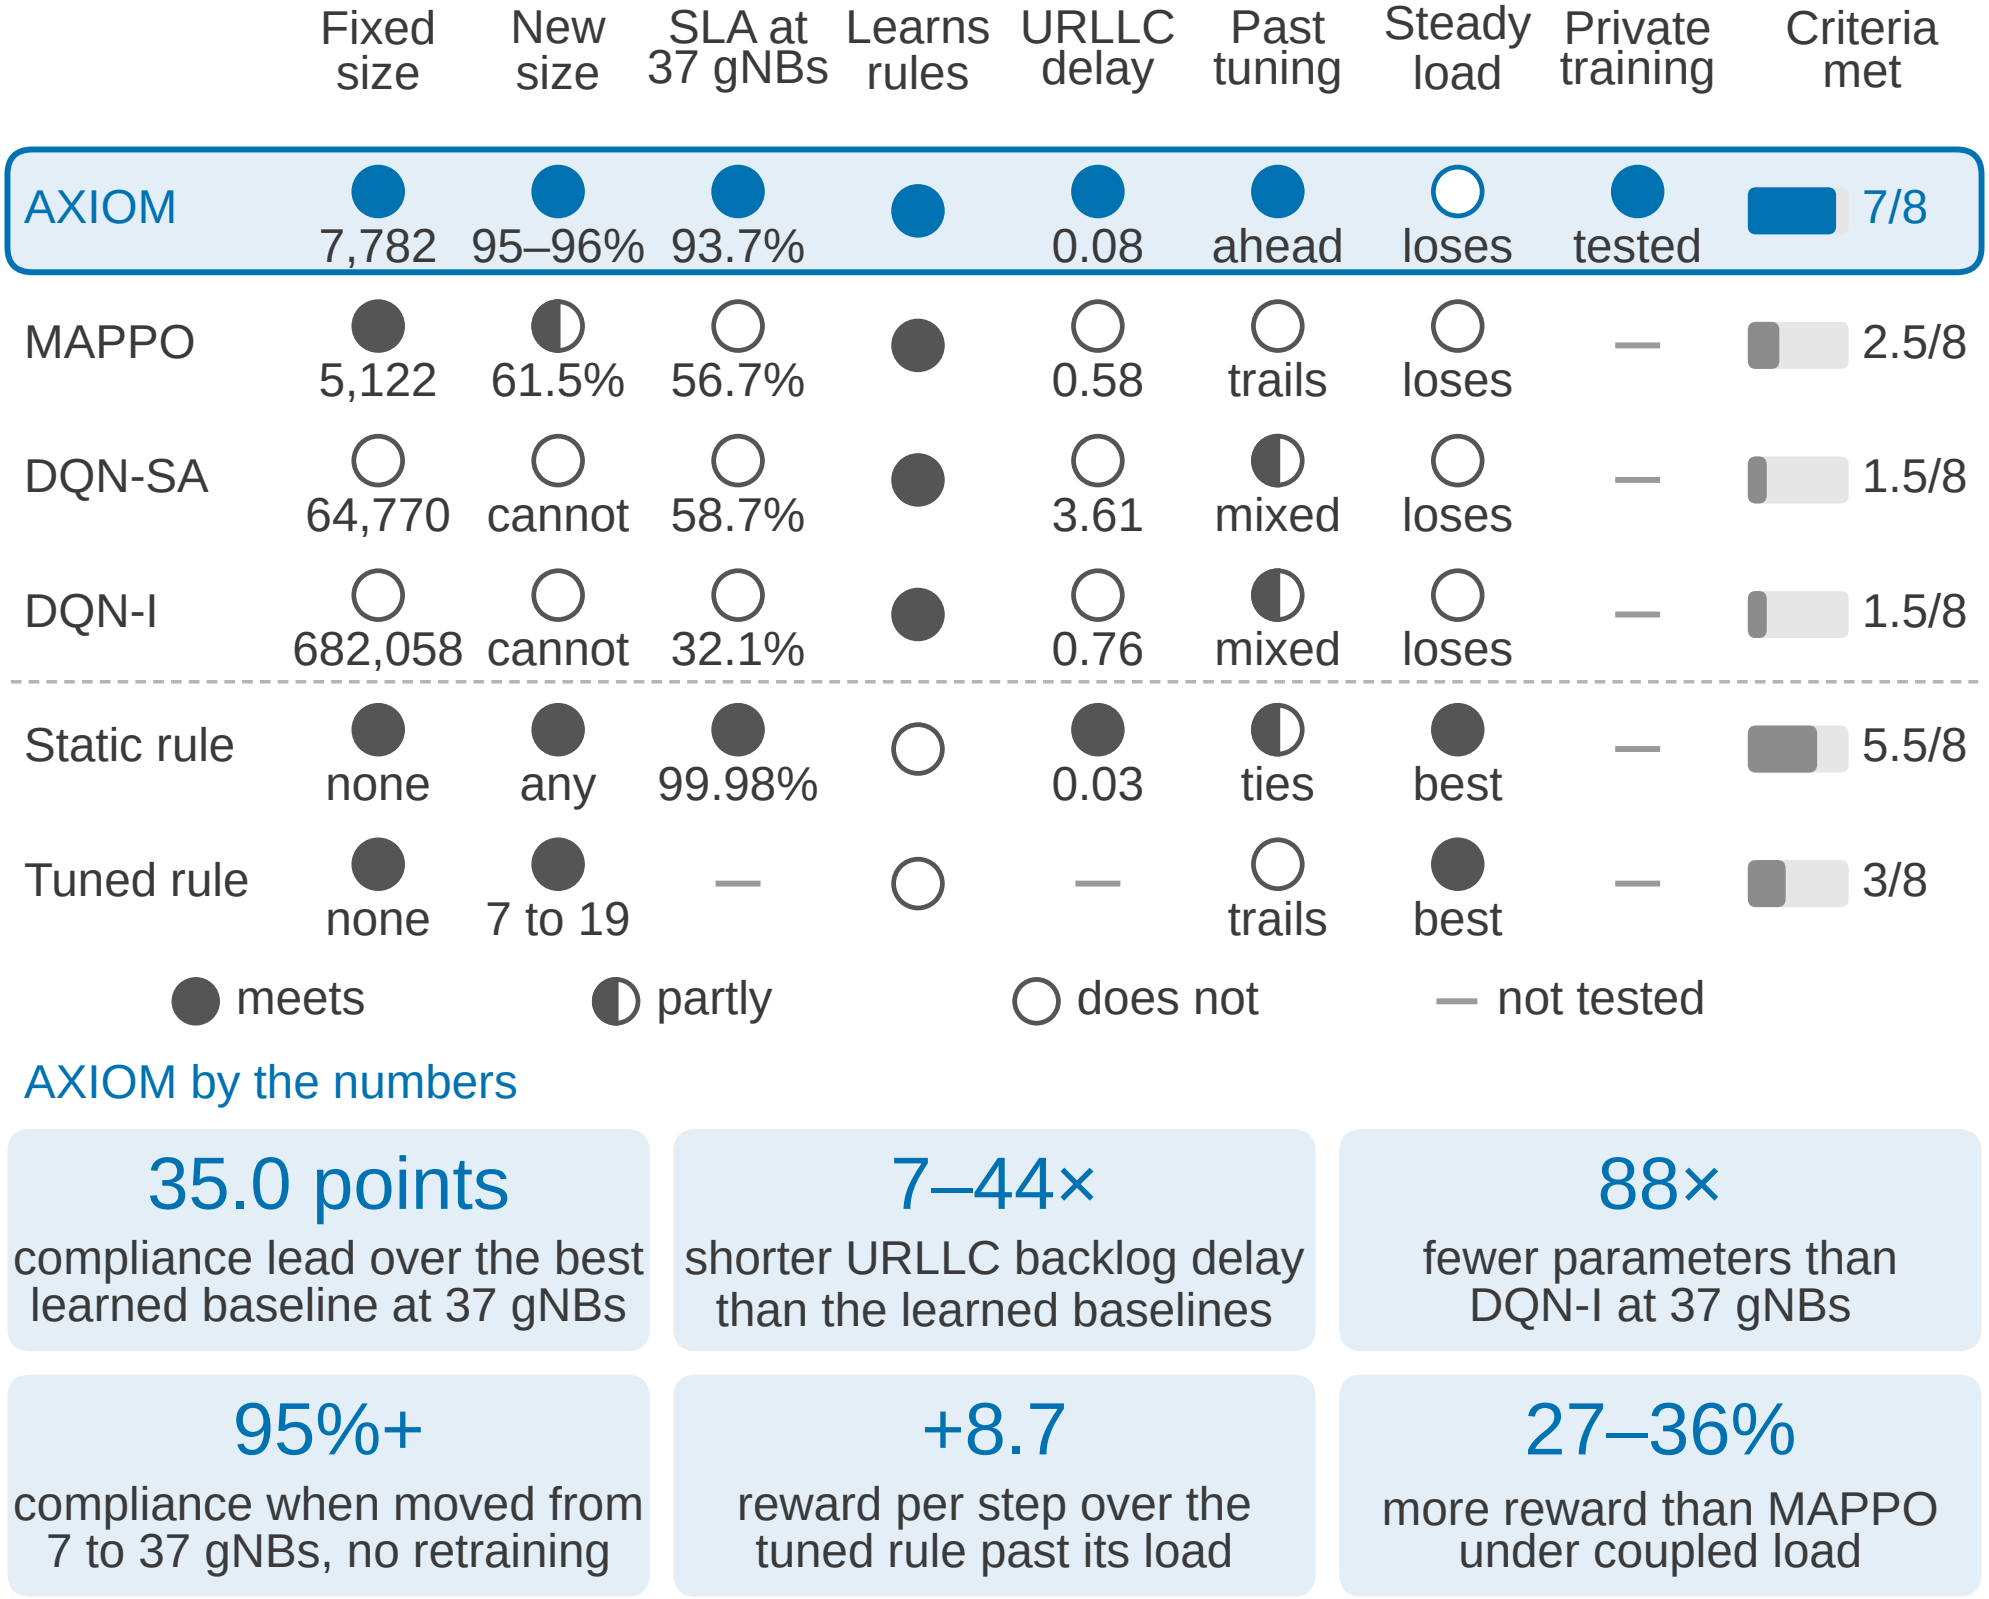

AXIOM by the numbers

35.0 points
compliance lead over the best learned baseline at 37 gNBs

7–44×
shorter URLLC backlog delay than the learned baselines

88×
fewer parameters than DQN-I at 37 gNBs

95%+
compliance when moved from 7 to 37 gNBs, no retraining

+8.7
reward per step over the tuned rule past its load

27–36%
more reward than MAPPO under coupled load