	Learned controllers				Rules	
	AXIOM	MAPPO	DQN-SA	DQN-I	Static rule	Tuned rule
Scale						
Fixed size	7,782	5,122	64,770	682,058	none	
New size	95–96%	61.5%	cannot		any	7 to 19
Service at 37 gNBs						
Priority SLA	93.7%	56.7%	58.7%	32.1%	99.98%	not tested
URLLC delay	0.08	0.58	3.61	0.76	0.03	not tested
Adaptation						
Learns its rules	yes				no	
Past tuning load	ahead	trails	mixed		ties	trails
Steady load	loses				best	
Privacy						
Private training	tested	not tested			not tested	
Criteria met	7/8	2.5/8	1.5/8	1.5/8	5.5/8	3/8

Cell colour: meets | partly | does not

AXIOM by the numbers

35.0 points
compliance lead over the best learned baseline at 37 gNBs

7–44×
shorter URLLC backlog delay than the learned baselines

88×
fewer parameters than DQN-I at 37 gNBs

95%+
compliance when moved from 7 to 37 gNBs, no retraining

+8.7
reward per step over the tuned rule past its load

27–36%
more reward than MAPPO under coupled load